



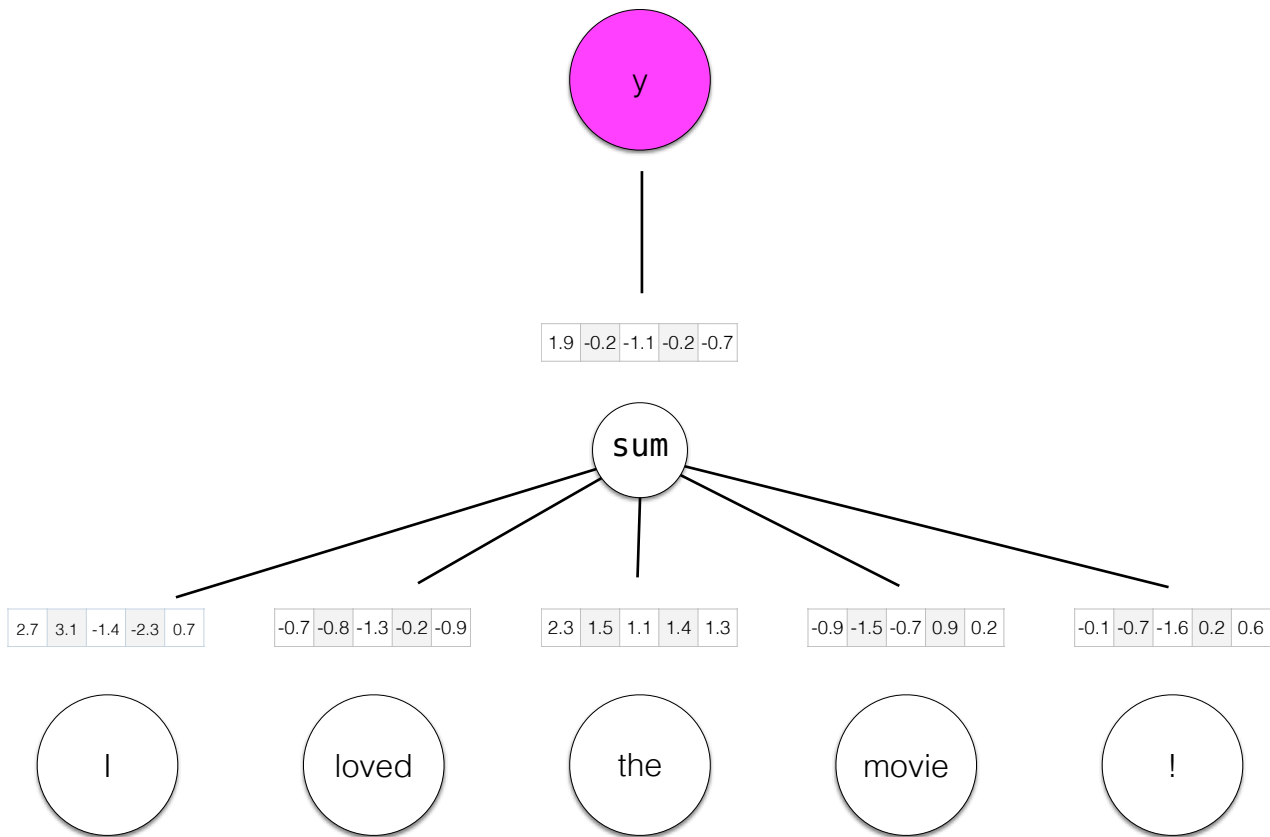
# Applied Natural Language Processing

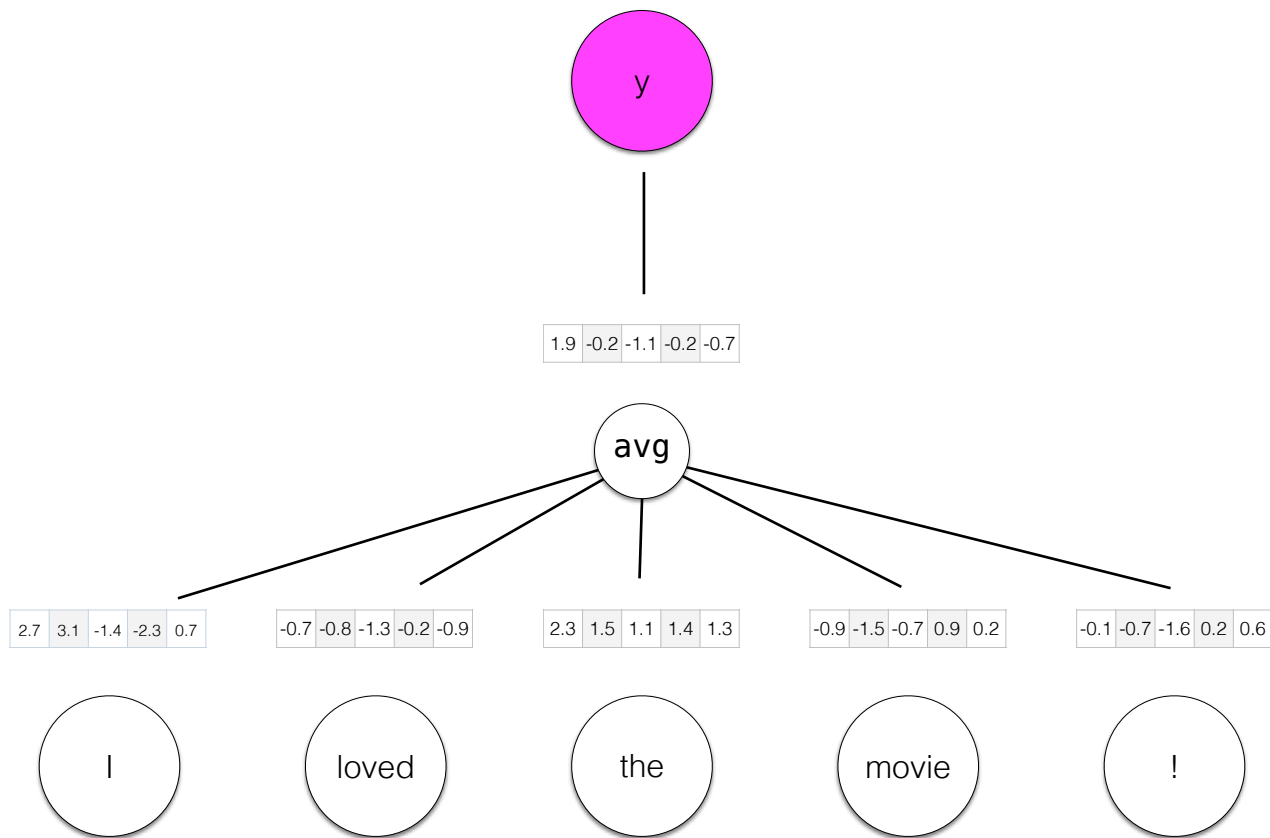
Info 256

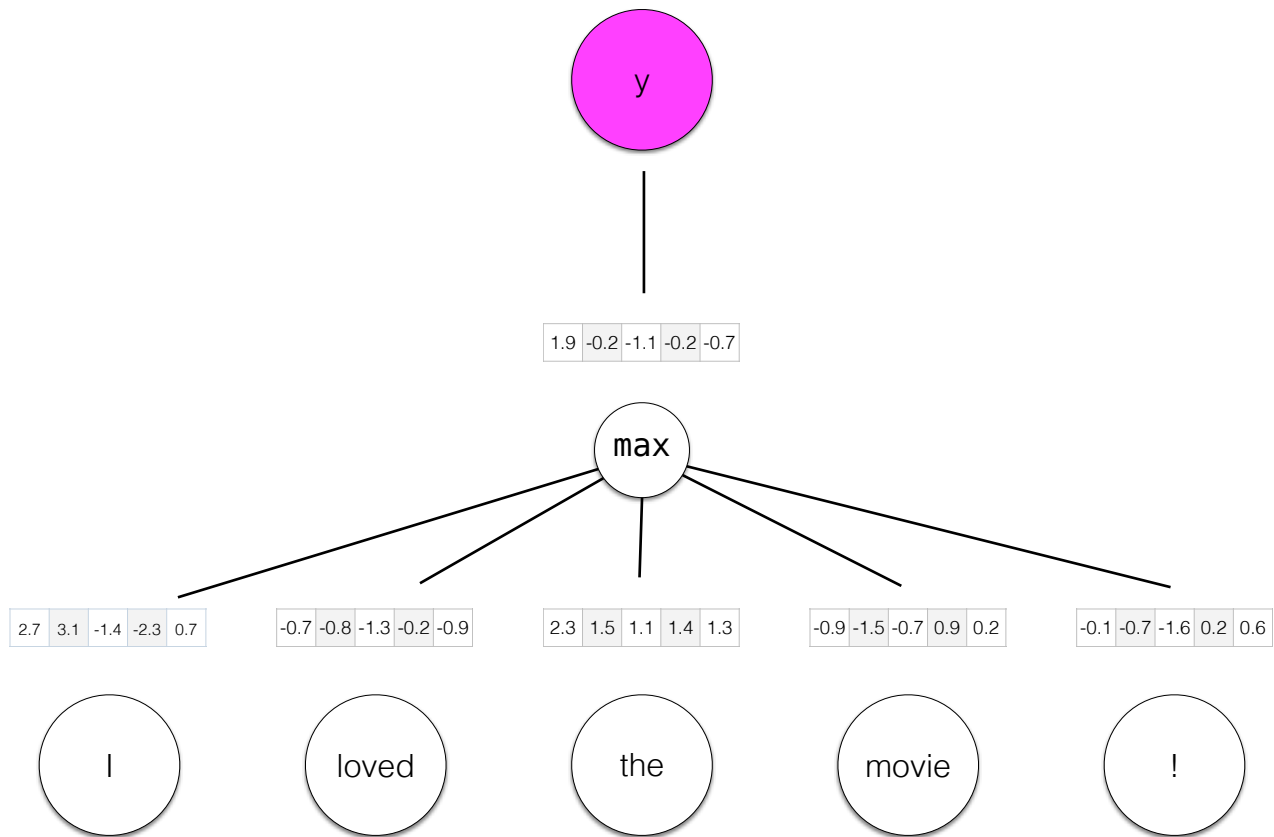
Lecture 15: Attention/BERT (Oct. 14, 2021)

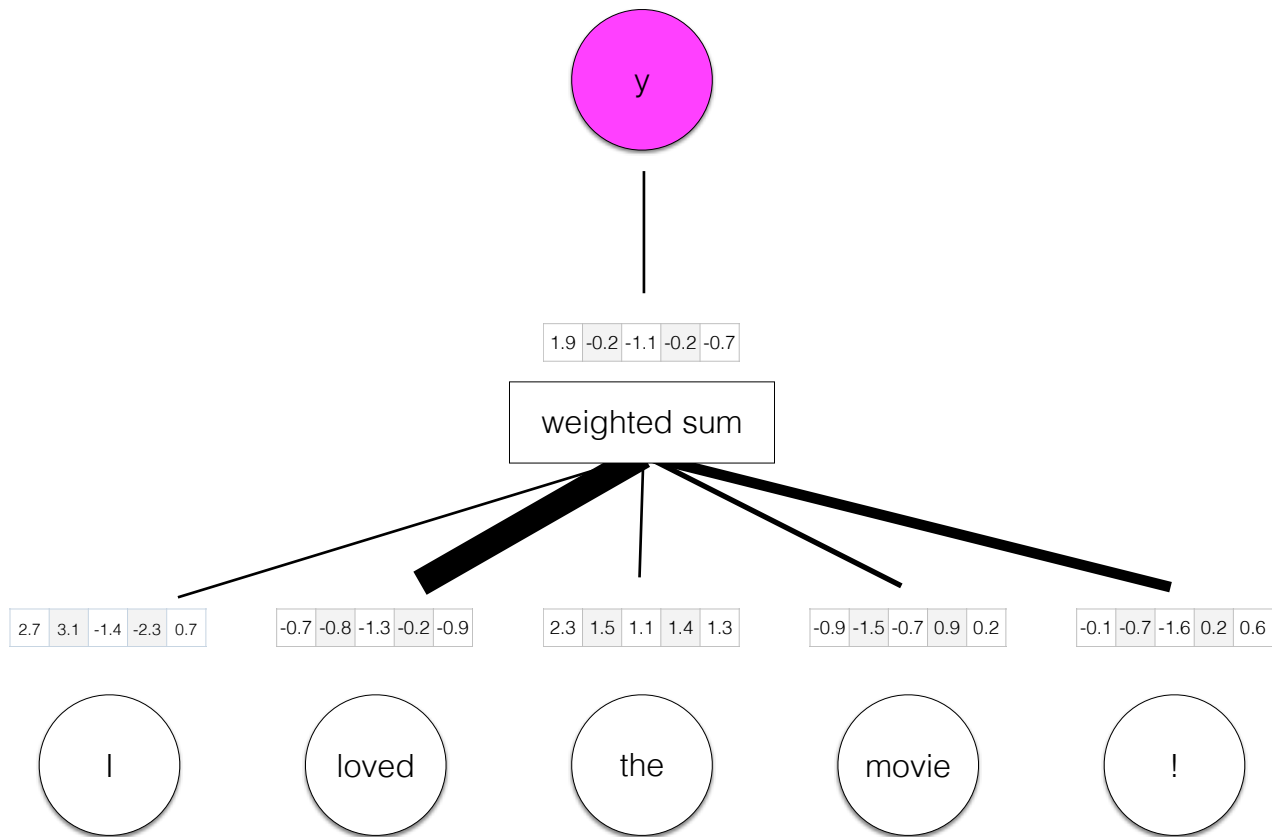
David Bamman, UC Berkeley

How do we use word embeddings for  
document classification?









# Attention


- Let's incorporate structure (and parameters) into a network that captures which elements in the input we should be **attending** to (and which we can ignore).

$$v \in \mathcal{R}^H$$

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----


Define  $v$  to be a vector to be learned; think of it as an “important word” vector. The dot product here measures how similar each input vector is to that “important word” vector

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----




$x_1$

-0.7	-0.8	-1.3	-0.2	-0.9
------	------	------	------	------




$x_2$

2.3	1.5	1.1	1.4	1.3
-----	-----	-----	-----	-----




$x_3$

-0.9	-1.5	-0.7	0.9	0.2
------	------	------	-----	-----



$x_4$

-0.1	-0.7	-1.6	0.2	0.6
------	------	------	-----	-----



$x_5$



$$v \in \mathcal{R}^H$$

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----

-3.4

2.4

-0.8


-1.2

1.7

$$r_1 = v^\top x_1$$

|

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----




$x_1$

$$r_2 = v^\top x_2$$

|

-0.7	-0.8	-1.3	-0.2	-0.9
------	------	------	------	------




$x_2$

$$r_3 = v^\top x_3$$

|

2.3	1.5	1.1	1.4	1.3
-----	-----	-----	-----	-----




$x_3$

$$r_4 = v^\top x_4$$

|

-0.9	-1.5	-0.7	0.9	0.2
------	------	------	-----	-----




$x_4$

$$r_5 = v^\top x_5$$

|

-0.1	-0.7	-1.6	0.2	0.6
------	------	------	-----	-----



$x_5$

Convert  $r$  into a vector of normalized weights that sum to 1.


$$a = \text{softmax}(r)$$

$a$	0	0.64	0.02	0.02	0.32
$r$	-3.4	2.4	-0.8	-1.2	1.7

$$r_1 = v^\top x_1$$

|

2.7	3.1	-1.4	-2.3	0.7
-----	-----	------	------	-----




$x_1$

$$r_2 = v^\top x_2$$

|

-0.7	-0.8	-1.3	-0.2	-0.9
------	------	------	------	------




$x_2$

$$r_3 = v^\top x_3$$

|

2.3	1.5	1.1	1.4	1.3
-----	-----	-----	-----	-----




$x_3$

$$r_4 = v^\top x_4$$

|

-0.9	-1.5	-0.7	0.9	0.2
------	------	------	-----	-----




$x_4$

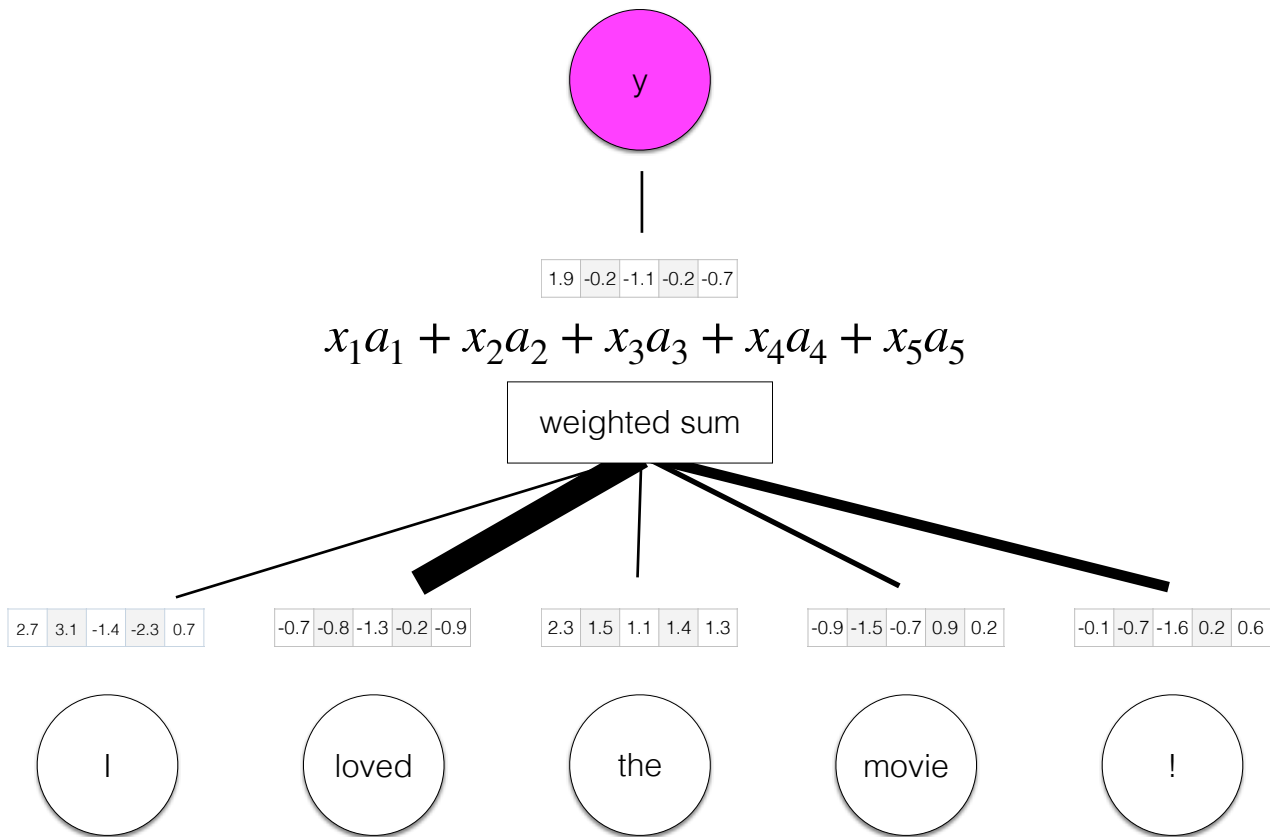
$$r_5 = v^\top x_5$$

|

-0.1	-0.7	-1.6	0.2	0.6
------	------	------	-----	-----



$x_5$



# Attention

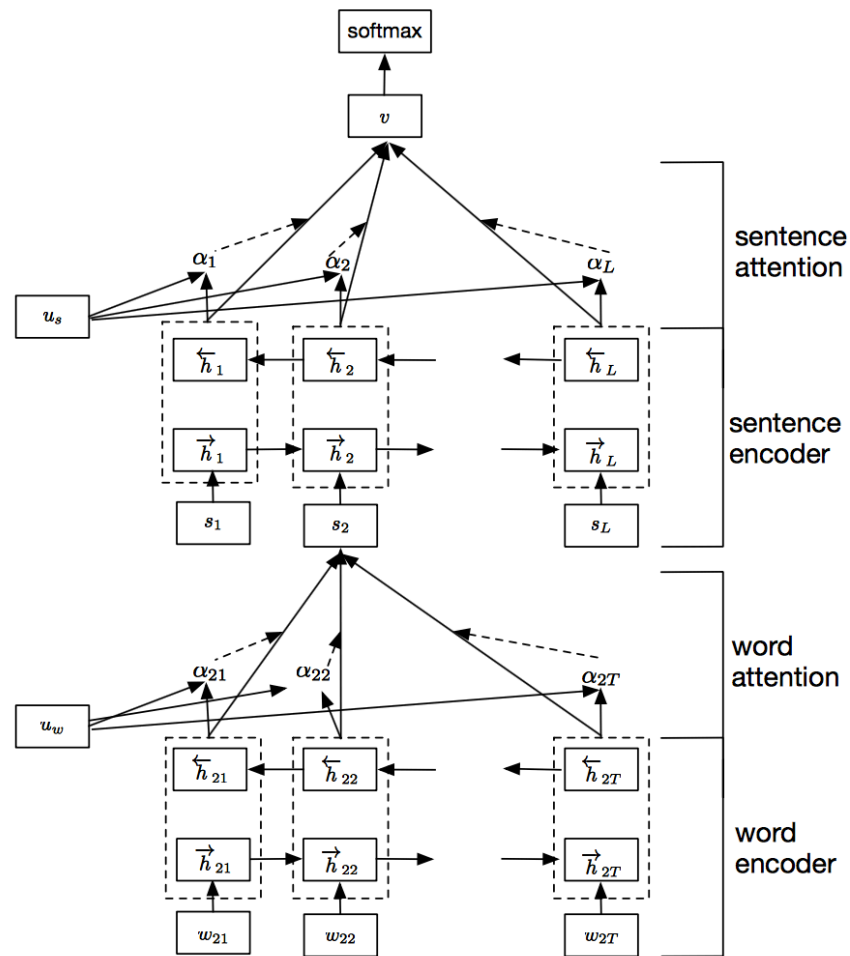
- Lots of variations on attention:
  - Linear transformation of  $x$  into before dotting with  $v$
  - Non-linearities after each operation.
  - “Multi-head attention”: multiple  $v$  vectors to capture different phenomena that can be attended to in the input.
  - Hierarchical attention (sentence representation with attention over words + document representation with attention over sentences).

attention over sentences

bidirectional GRU over  
sentence representations


attention over words

bidirectional GRU over  
word representations



# Attention

- Attention gives us a normalized weight for every token in a sequence that tells us how important that word was for the prediction
- This can be useful for visualization



# BERT

- Transformer-based model (Vaswani et al. 2017) to predict masked word using bidirectional context + next sentence prediction.
- Generates multiple layers of representations for each token sensitive to its context of use.

Each token in the input starts out represented by token and position embeddings

-0.2	1	0.1	-0.8	-1.1
$e_{1,1}$				

The

0.3	0.3	-1.7	0.7	-1.1
$e_{1,2}$				

dog

1.6	-0.3	-0.9	-0.7	0.2
$e_{1,3}$				

barked



The value for time step  $j$  at layer  $i$  is the result of attention over all time steps in the previous layer  $i-1$

-0.7	-1.3	0.4	-0.4	-0.7
$e_{2,1}$				

-0.2	1	0.1	-0.8	-1.1
$e_{1,1}$				





The

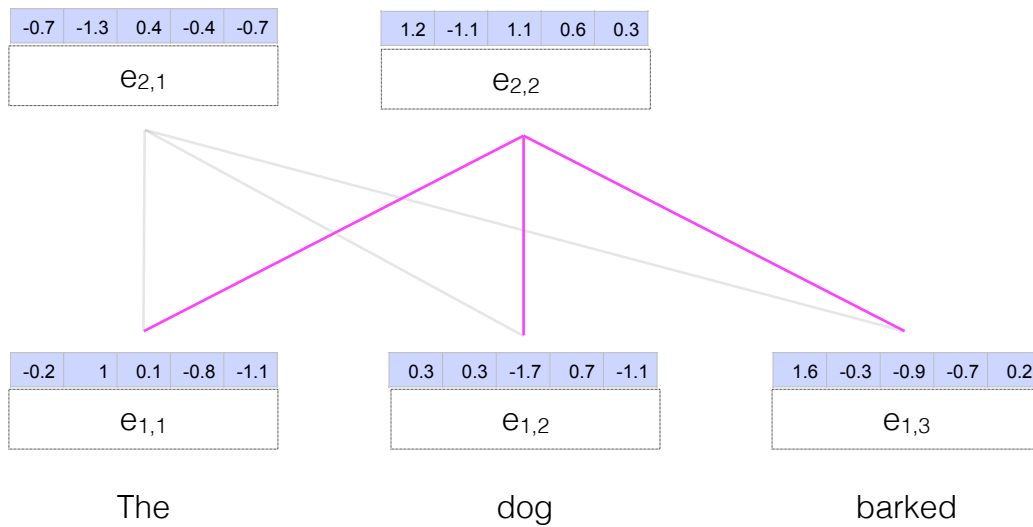
0.3	0.3	-1.7	0.7	-1.1
$e_{1,2}$				

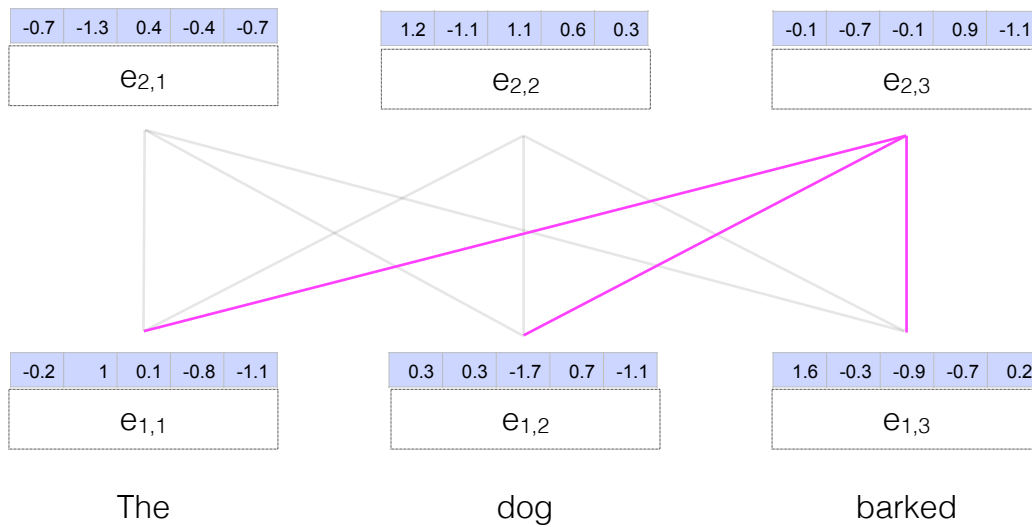
dog

1.6	-0.3	-0.9	-0.7	0.2
$e_{1,3}$				

barked







-0.2	0.3	2.1	1.2	0.6
$e_{3,1}$				

-0.7	-1.3	0.4	-0.4	-0.7
$e_{2,1}$				

1.2	-1.1	1.1	0.6	0.3
$e_{2,2}$				

-0.1	-0.7	-0.1	0.9	-1.1
$e_{2,3}$				

-0.2	1	0.1	-0.8	-1.1
$e_{1,1}$				


0.3	0.3	-1.7	0.7	-1.1
$e_{1,2}$				

1.6	-0.3	-0.9	-0.7	0.2
$e_{1,3}$				

The

dog

barked



-0.2	0.3	2.1	1.2	0.6
$e_{3,1}$				

-1.8	-0.2	-2.4	-0.2	-0.1
$e_{3,2}$				

-0.7	-1.3	0.4	-0.4	-0.7
$e_{2,1}$				

1.2	-1.1	1.1	0.6	0.3
$e_{2,2}$				

-0.1	-0.7	-0.1	0.9	-1.1
$e_{2,3}$				

-0.2	1	0.1	-0.8	-1.1
$e_{1,1}$				


0.3	0.3	-1.7	0.7	-1.1
$e_{1,2}$				

1.6	-0.3	-0.9	-0.7	0.2
$e_{1,3}$				

The

dog


barked



-0.2	0.3	2.1	1.2	0.6
$e_{3,1}$				

-1.8	-0.2	-2.4	-0.2	-0.1
$e_{3,2}$				

-0.9	-1.5	-0.7	0.9	0.2
$e_{3,3}$				



-0.7	-1.3	0.4	-0.4	-0.7
$e_{2,1}$				

1.2	-1.1	1.1	0.6	0.3
$e_{2,2}$				

-0.1	-0.7	-0.1	0.9	-1.1
$e_{2,3}$				



-0.2	1	0.1	-0.8	-1.1
$e_{1,1}$				

0.3	0.3	-1.7	0.7	-1.1
$e_{1,2}$				

1.6	-0.3	-0.9	-0.7	0.2
$e_{1,3}$				

The

dog

barked

At the end of this process, we have one  
representation for each layer for each token

-0.2	0.3	2.1	1.2	0.6
$e_{3,1}$				

-1.8	-0.2	-2.4	-0.2	-0.1
$e_{3,2}$				

-0.9	-1.5	-0.7	0.9	0.2
$e_{3,3}$				

-0.7	-1.3	0.4	-0.4	-0.7
$e_{2,1}$				

1.2	-1.1	1.1	0.6	0.3
$e_{2,2}$				

-0.1	-0.7	-0.1	0.9	-1.1
$e_{2,3}$				

-0.2	1	0.1	-0.8	-1.1
$e_{1,1}$				

0.3	0.3	-1.7	0.7	-1.1
$e_{1,2}$				

1.6	-0.3	-0.9	-0.7	0.2
$e_{1,3}$				

The

dog

barked

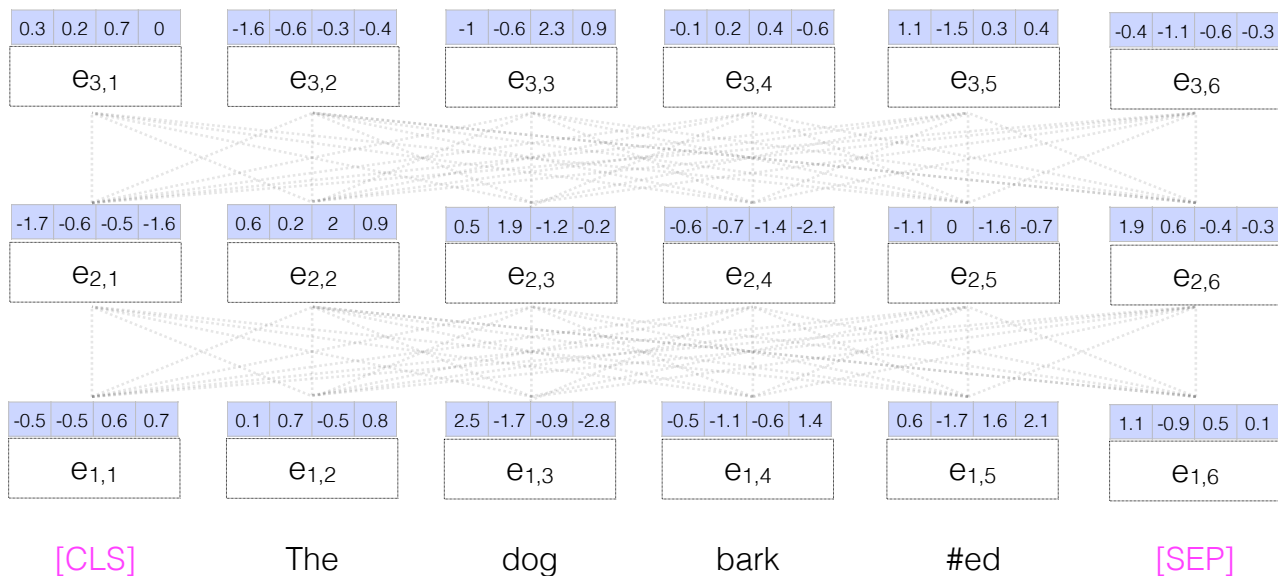


# WordPiece

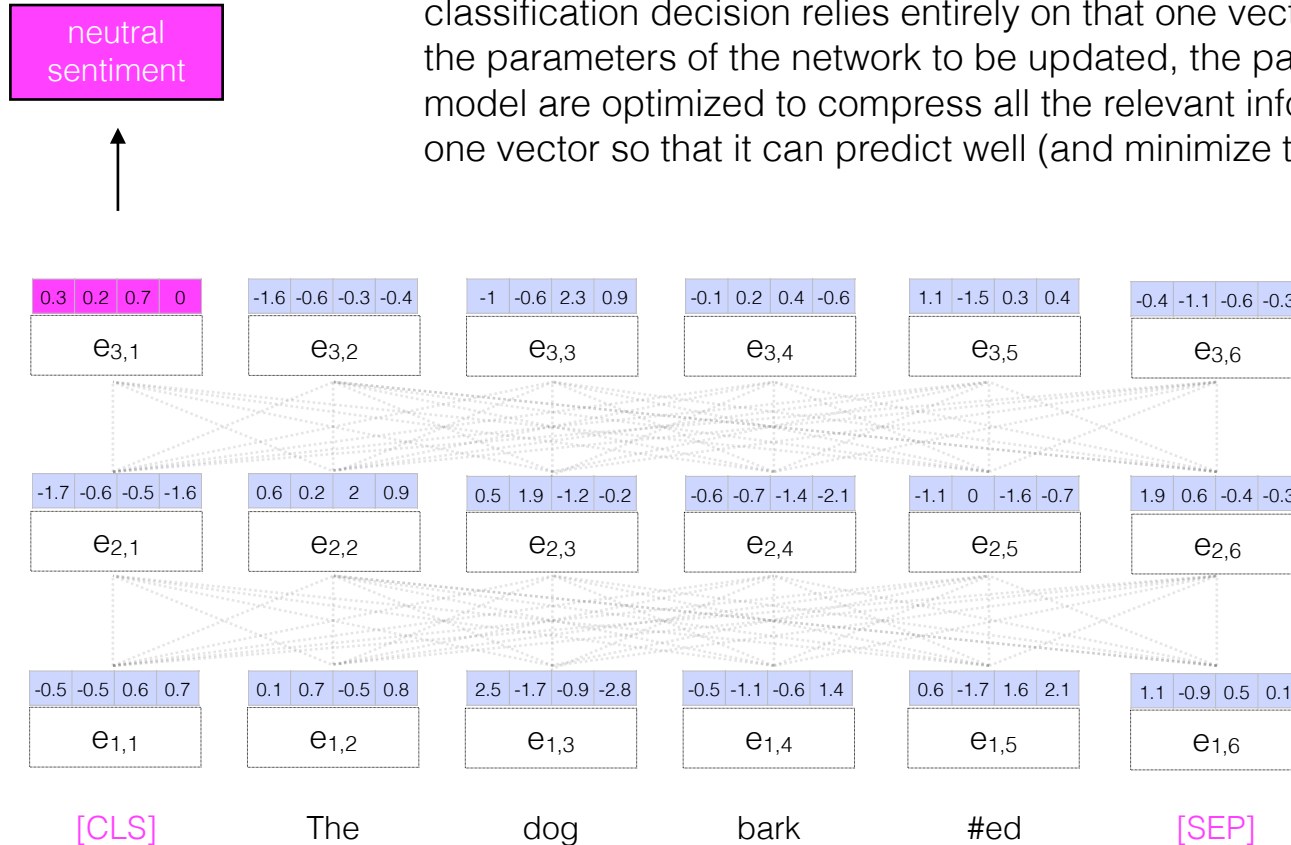
- BERT uses WordPiece tokenization, which segments some morphological structure of tokens
- Vocabulary size: 30,000

The	The
dog	dog
barked	bark #ed

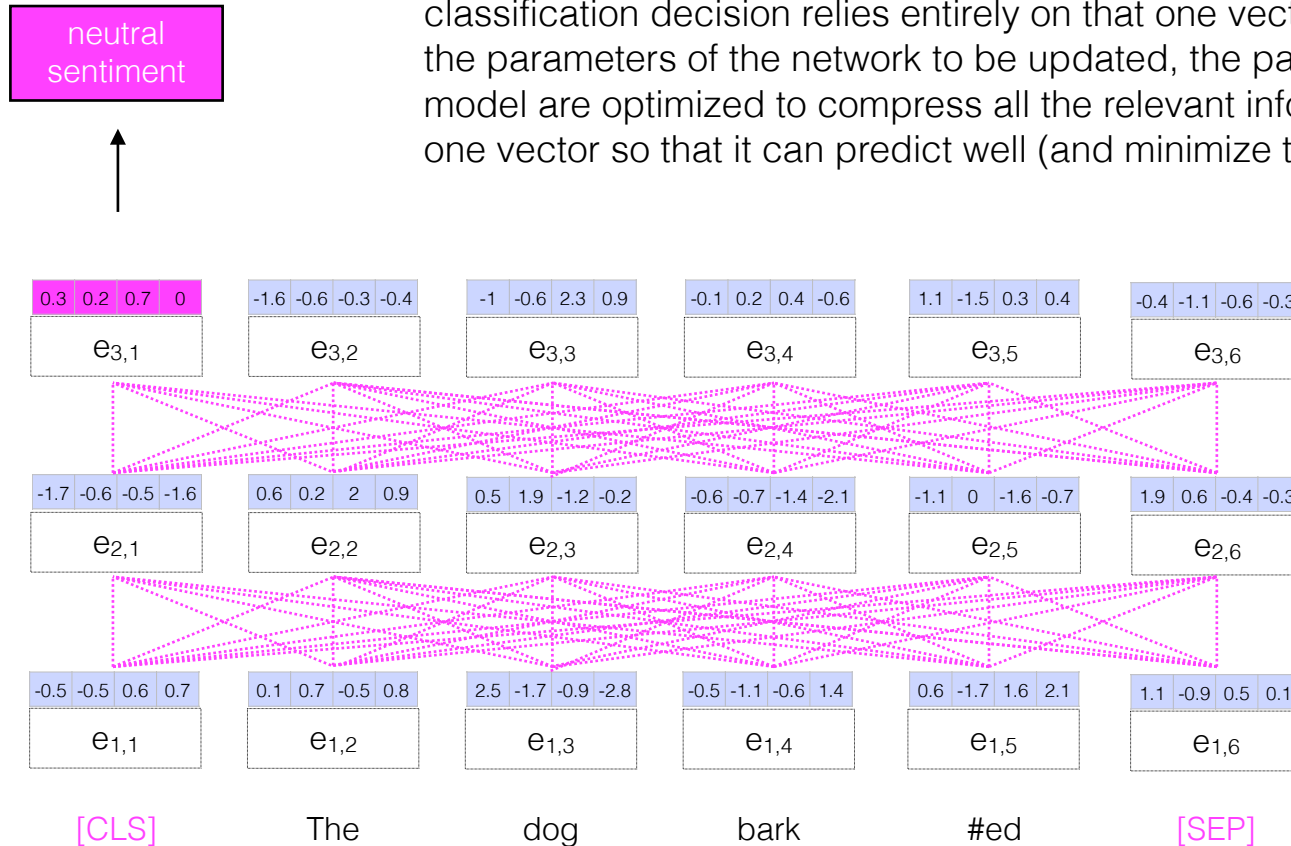
- BERT also encodes each sentence by appending a special token to the beginning ([CLS]) and end ([SEP]) of each sequence.
- This helps provides a single token that can be optimized to represent the entire sequence (e.g., for document classification)



- We can represent the entire document with this *one* [CLS] vector
- Why does this work? When we design our network so that a classification decision relies entirely on that one vector and allow all the parameters of the network to be updated, the parameters of the model are optimized to compress all the relevant information into that one vector so that it can predict well (and minimize the loss).



- We can represent the entire document with this *one* [CLS] vector
- Why does this work? When we design our network so that a classification decision relies entirely on that one vector and allow all the parameters of the network to be updated, the parameters of the model are optimized to compress all the relevant information into that one vector so that it can predict well (and minimize the loss).



# BERT

- Deep layers (12 for BERT base, 24 for BERT large)
- Large representation sizes (768 per layer)
- Pretrained on English Wikipedia (2.5B words) and BooksCorpus (800M words)

# BERT

	H=128	H=256	H=512	H=768
L=2	2/128 (BERT-Tiny)	2/256	2/512	2/768
L=4	4/128	4/256 (BERT-Mini)	4/512 (BERT-Small)	4/768
L=6	6/128	6/256	6/512	6/768
L=8	8/128	8/256	8/512 (BERT-Medium)	8/768
L=10	10/128	10/256	10/512	10/768
L=12	12/128	12/256	12/512	12/768 (BERT-Base)

<https://github.com/google-research/bert>



v4.11.3 ▼

🏠 transformers



52,449

## GET STARTED

Quick tour

Installation

Philosophy

Glossary

## USING 🧑🏻 TRANSFORMERS

Summary of the tasks

Summary of the models

Preprocessing data

🔥 SIGN IN

🚀 MODELS

💬 FORUM

[Docs](#) » Pretrained models

[View page source](#)

# Pretrained models 🔗

Here is a partial list of some of the available pretrained models together with a short presentation of each model.

For the full list, refer to <https://huggingface.co/models>.

Architecture	Model id	Details of the model
	<code>bert-base-uncased</code>	12-layer, 768-hidden, 12-heads, 110M parameters. Trained on lower-cased English text.
	<code>bert-large-uncased</code>	24-layer, 1024-hidden, 16-heads, 336M parameters. Trained on lower-cased English text.

[https://huggingface.co/transformers/pretrained\\_models.html](https://huggingface.co/transformers/pretrained_models.html)



## Lost in (language-specific) BERT models? We are here to help!

We currently have indexed **31** BERT-based models, **19** Languages and **28** Tasks.

We have a total of **178** entries in this table; we also show **Multilingual Bert (mBERT)** results if available! (see our [paper](#))

Curious which BERT model is the best for named entity recognition in Italian 🇮🇹? Just type *"Italian NER"* in the search bar!


Show  entries

Search:

Language ↕	Model ↕	NLP Task ↕	Dataset ↕	Dataset-Domain ↕	Measure ↕	Performance ↕	mBERT ↕	Difference with mBERT ↕	Source ↕
Arabic 🇸🇦	Arabert v1	SA	AJGT	twitter	Accuracy	93.8	83.6	10.2	<a href="#">🔗</a> <a href="#">🔄</a>
Arabic 🇸🇦	Arabert v1	SA	HARD	hotel reviews	Accuracy	96.1	95.7	0.4	<a href="#">🔗</a> <a href="#">🔄</a>
Arabic 🇸🇦	Arabert v1	SA	ASTD	twitter	Accuracy	92.6	80.1	12.5	<a href="#">🔗</a> <a href="#">🔄</a>
Arabic 🇸🇦	Arabert v1	SA	ArSenTD-Lev	twitter	Accuracy	59.4	51.0	8.4	<a href="#">🔗</a> <a href="#">🔄</a>



# Progress — Coreference resolution



# Bertology

- Hewitt et al. 2019
- Tenney et al. 2019
- McCoy et al. 2019
- Liu et al. 2019
- Clark et al. 2019
- Goldberg 2019
- Michel et al. 2019


# Code

Pre-trained models for BERT, Transformer-XL, ALBERT, RoBERTa, DistilBERT, GPT-2, etc. for English, French, “Multilingual”

<https://huggingface.co>

# Probing

- Even though BERT is mainly trained on a language modeling objective, it learns a lot about the structure of language — even without direct training data for specific linguistic tasks.
- Probing experiments uncover what—and where (in what layers)—pretrained BERT encodes this information.



# Activity

`9.neural/BERTClassification`

- Explore BERT for document classification using [Google Colab](#)